



Aggregated Optuna archive: 79 real-hardware trials, 61 valid runs, 87.9 min closed-loop time



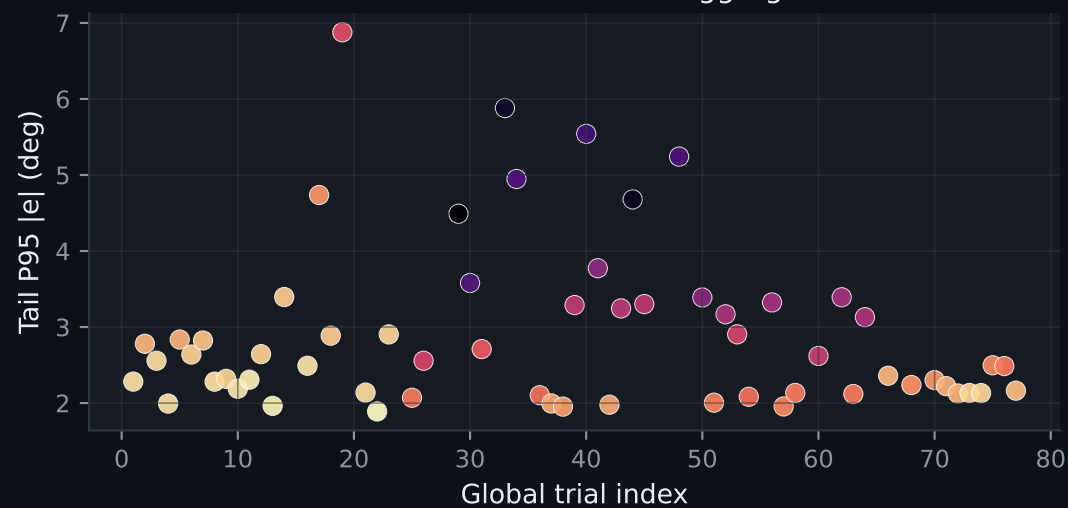
Accuracy versus bounded adaptation



Tail error versus saturation (marker size = activity)



Late-window residuals across the aggregated archive



Dominant parameter-score associations

